



# Architektur und Implementierung von Datenbanksystemen

## Task 8

Team 3 - Gründlinger Diana, Huber Marcel, Klotz Thomas, Targa Aaron, Thalmann Matthias

# PostgreSQL cost estimation

- Each parameter can be customized
- ANALYZE collects statistics about contents of tables → *pg\_statistic*  
→ statistics used by query planner (as partially seen in lecture 08: Optimization 2)
- ANALYZE is automatically done by the *autovacuum daemon*
- *pg\_statistic*:
  - entry for each column
  - fraction of null values, avg. stored bytes, amount of distinct values, ...

# Cost estimation parameters (1)

- Arbitrary scale
- Relative values
- Default: relative to *seq\_page\_cost*

| Parameter            | Default value | Description  |
|----------------------|---------------|--|
| seq_page_cost        | 1.0           | Cost of a disk page fetch (sequentially)                 |
| random_page_cost     | 4.0           | Cost of a disk page fetch (non-sequentially)             |
| cpu_tuple_cost       | 0.01          | Cost of processing each row during a query               |
| cpu_index_tuple_cost | 0.005         | Cost of processing each index entry during an index scan |
| cpu_operator_cost    | 0.0025        | Cost of processing each operator executed during a query |

## Cost estimation parameters (2)

| Parameter                    | Default value | Description  |
|------------------------------|---------------|--|
| parallel_setup_cost          | 1000          | Cost of launching parallel worker processes  |
| parallel_tuple_cost          | 0.1           | Cost of transferring a tuple from one process to another one                               |
| min_parallel_table_scan_size | 8MB           | Min. amount of <u>table data</u> needed to be scanned for a parallel scan to be considered |
| min_parallel_index_scan_size | 512KB         | Min. amount of <u>index data</u> needed to be scanned for a parallel scan to be considered |
| effective_cache_size         | 4GB           | Assumed amount of disk cache avail. to a single query                                      |
| jit_above_cost               | 100,000       | Cost, above which JIT compilation is activated   |
| jit_inline_above_cost        | 500,000       | Cost, above which JIT comp. attempts inlining  |
| jit_optimize_above_cost      | 500,000       | Cost, above which JIT comp. applies expensive optimization                                 |

# Cost estimation – Sequential scan

- $(\text{disk\_page\_reads} * \text{seq\_page\_cost}) + (\text{rows\_scanned} * \text{cpu\_tuple\_cost})$
- Filtering (*WHERE*-clause) does only increase the cost, since each row has to be scanned *and* evaluated  
→  $+ (\text{rows\_scanned} * \text{cpu\_operator\_cost})$

## Example 1:

*tenk1* has 358 disk pages and 10,000 rows

```
EXPLAIN SELECT * FROM tenk1;
```

QUERY PLAN

```
-----  
Seq Scan on tenk1 (cost=0.00..458.00 rows=10000)
```

$$\rightarrow (358 * 1.0) + (10000 * 0.01) = 458.0$$

## Example 2:

*tenk1* has 358 disk pages and 10,000 rows

```
EXPLAIN SELECT * FROM tenk1 WHERE unique1 < 7000;
```

QUERY PLAN

```
-----  
Seq Scan on tenk1 (cost=0.00..483.00 rows=7001 width=24)  
Filter: (unique1 < 7000)
```

$$\begin{aligned} \rightarrow & (358 * 1.0) + (10000 * 0.01) \\ & + (10000 * 0.0025) = 483.0 \end{aligned}$$

# Cost estimation – Index scan (1)

- Very complex method in *costsize.c*
- Heavily dependent on the implementation of the index (e.g. B-Tree index)
- *amcostestimate* – function-interface; implementation specific for each type of index
- Estimates the total cost of the index (and other costs)
  - depending on type of index and index statistics
  - uses the *genericcostestimate* function  
(takes for example caching into account)

```
EXPLAIN SELECT * FROM tenk1 WHERE unique1 = 42;
```

QUERY PLAN

---

Index Scan using tenk1\_unique1 on tenk1 (cost=0.29..8.30 rows=1 w  
Index Cond: (unique1 = 42)

# Cost estimation – Index scan (2) – Example for a B-Tree index

```
WITH costs(idx_cost, tbl_cost) AS (  
  SELECT  
    (  
      SELECT round(  
        current_setting('random_page_cost')::real * pages +  
        current_setting('cpu_index_tuple_cost')::real * tuples +  
        current_setting('cpu_operator_cost')::real * tuples  
      )  
    FROM (  
      SELECT relpages * 0.0630 AS pages, reltuples * 0.0630 AS tuples  
      FROM pg_class WHERE relname = 'bookings_pkey'  
    ) c  
    ),  
    (  
      SELECT round(  
        current_setting('seq_page_cost')::real * pages +  
        current_setting('cpu_tuple_cost')::real * tuples  
      )  
    FROM (  
      SELECT relpages * 0.0630 AS pages, reltuples * 0.0630 AS tuples  
      FROM pg_class WHERE relname = 'bookings'  
    ) c  
    )  
  )  
  SELECT idx_cost, tbl_cost, idx_cost + tbl_cost AS total FROM costs;
```

Amount of rows

SELECT round(132999::numeric/reltuples::numeric, 4)  
FROM pg\_class WHERE relname = 'bookings';

round  
-----  
0.0630  
(1 row)

Source:

<https://postgrespro.com/blog/pgsql/5969493>

# References

- <https://github.com/postgres/postgres/blob/master/src/backend/optimizer/path/costsize.c>
- <http://morningcoffee.io/the-postgresql-query-cost-model.html>
- <https://www.postgresql.org/docs/current/runtime-config-query.html>
- <https://www.postgresql.org/docs/current/using-explain.html>
- <https://www.postgresql.org/docs/current/index-cost-estimation.html>
- <https://www.postgresql.org/docs/current/sql-analyze.html>
- <https://www.postgresql.org/docs/current/routine-vacuuming.html>
- [https://git.uibk.ac.at/informatik/dbis/dbis-teaching/archimpl-course-material-2022/-/blob/main/slides/08\\_optimization2.pdf](https://git.uibk.ac.at/informatik/dbis/dbis-teaching/archimpl-course-material-2022/-/blob/main/slides/08_optimization2.pdf)
- <https://postgrespro.com/blog/pgsql/5969493>
- <https://stackoverflow.com/a/57279653/11028838>



